GRIT: Enhancing Multi-GPU Performance with Fine-Grained Dynamic Page Placement

Yueqi Wang*[1], ***Bingyao Li****[1], Aamer Jaleel[2], Jun Yang[1], Xulong Tang[1]

[1]University of Pittsburgh, [2]NVIDIA

Multi-GPU is Popular

Graph Processing

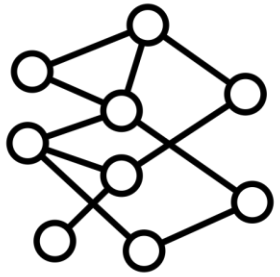

DNN

ChatGPT

Datacenter Workloads

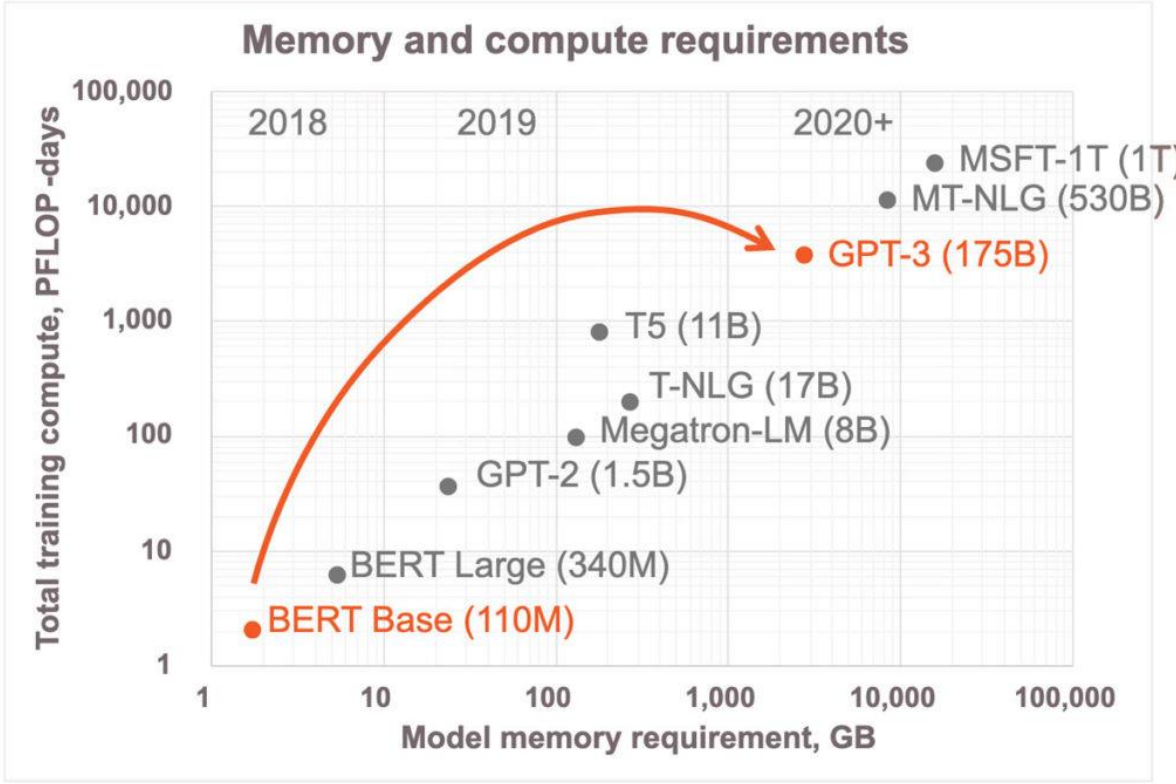

Ever-growing application complexity and input dataset sizes.

University of Pittsburgh

[1] https://www.cerebras.net/blog/harnessing-the-power-of-sparsity-for-large-gpt-ai-models

Multi-GPU is Popular

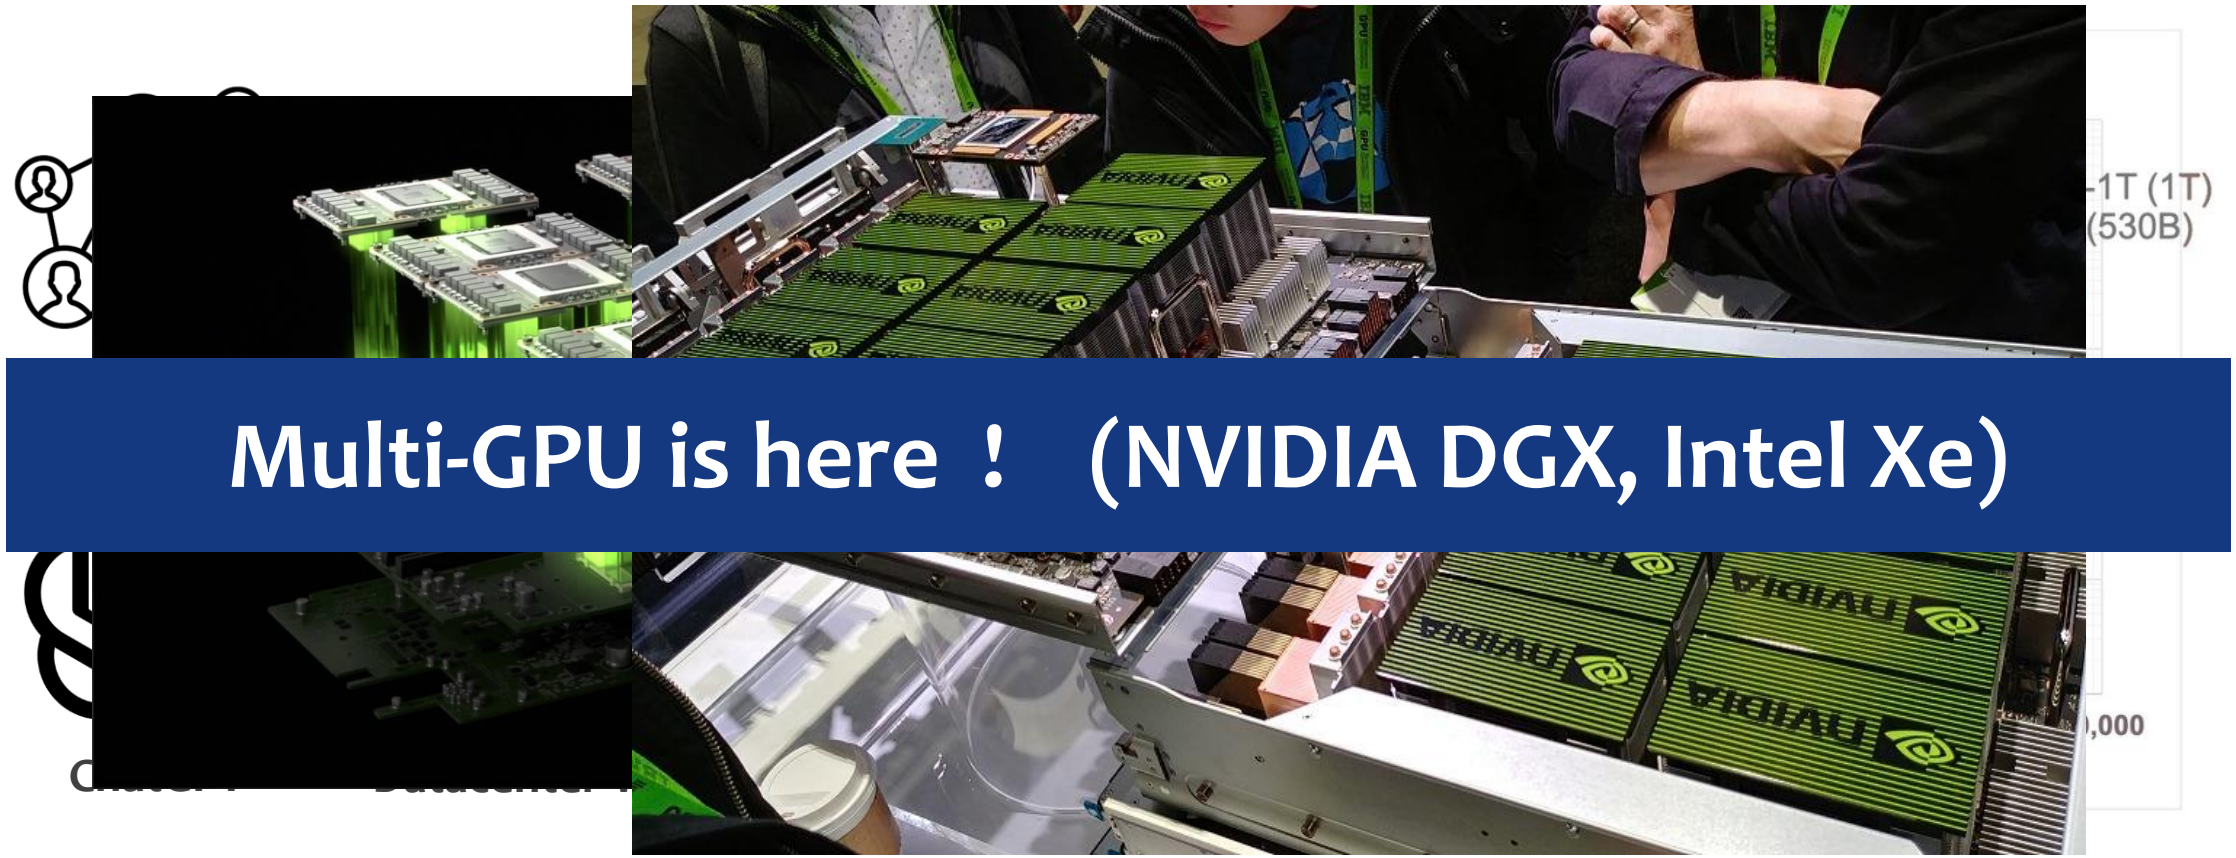

Multi-GPU is here ！ (NVIDIA DGX, Intel Xe)

Ever-growing application complexity and input dataset sizes.

[1] https://www.cerebras.net/blog/harnessing-the-power-of-sparsity-for-large-gpt-ai-models

UVM for Multi-GPU

- Growing trend of multi-GPUs leveraging **Unified Virtual Memory** (UVM)

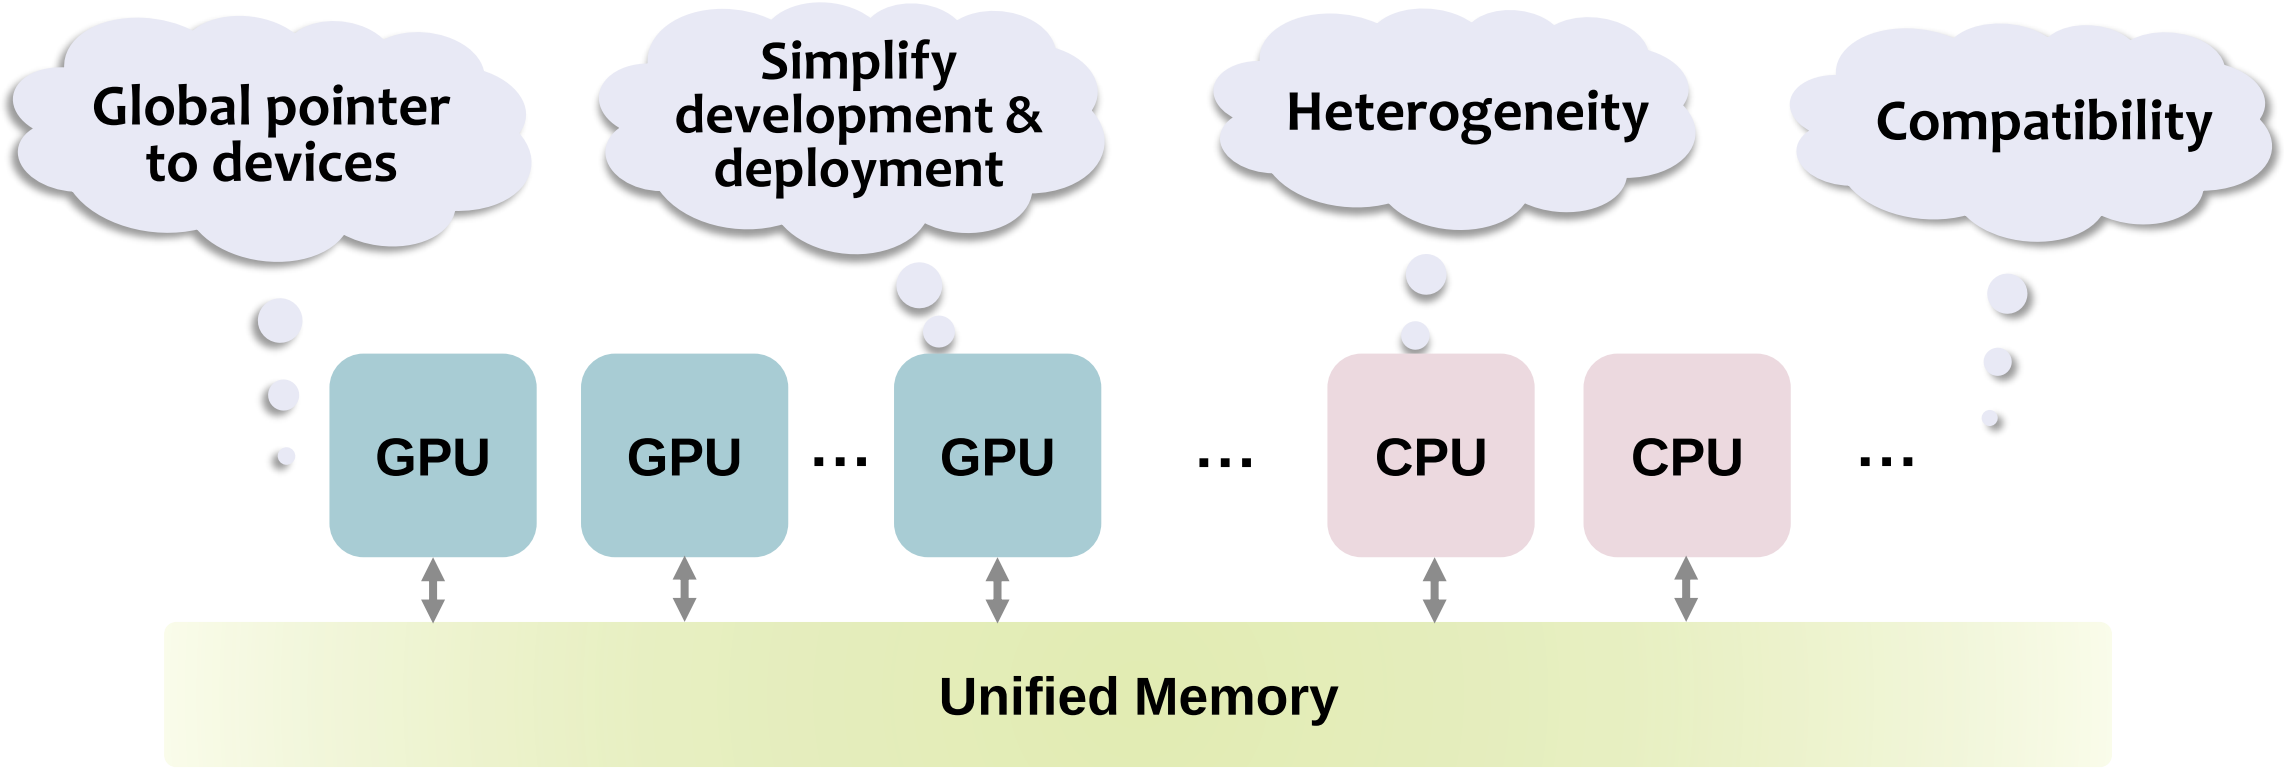

Multi-GPU Scalability

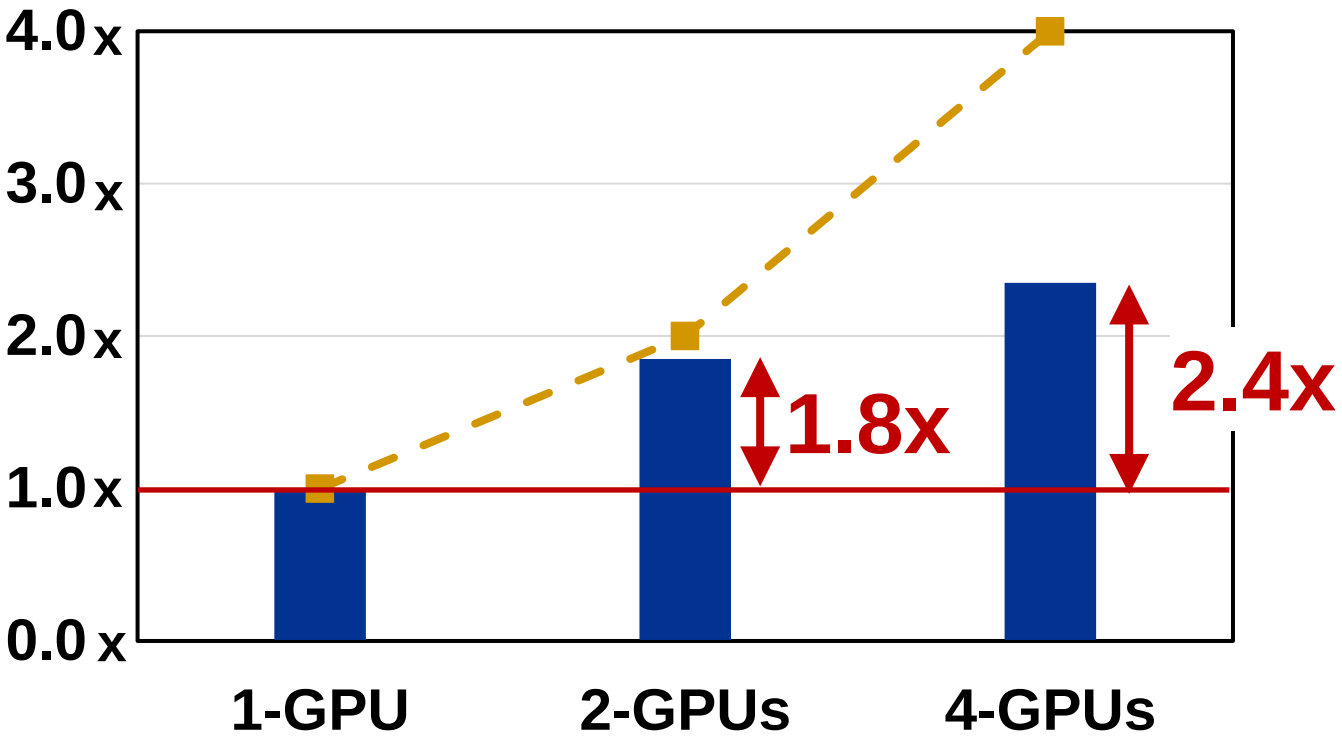

Why performance gap:

- **NUMA data access**
- **Data transfer**
- **Address translation**

……

[2] https://developer.nvidia.com/blog/easy-multi-gpu-deep-learning-digits-2/

Multi-GPU Page Placement Schemes

1. On-Touch Migration: Request, Migrate.

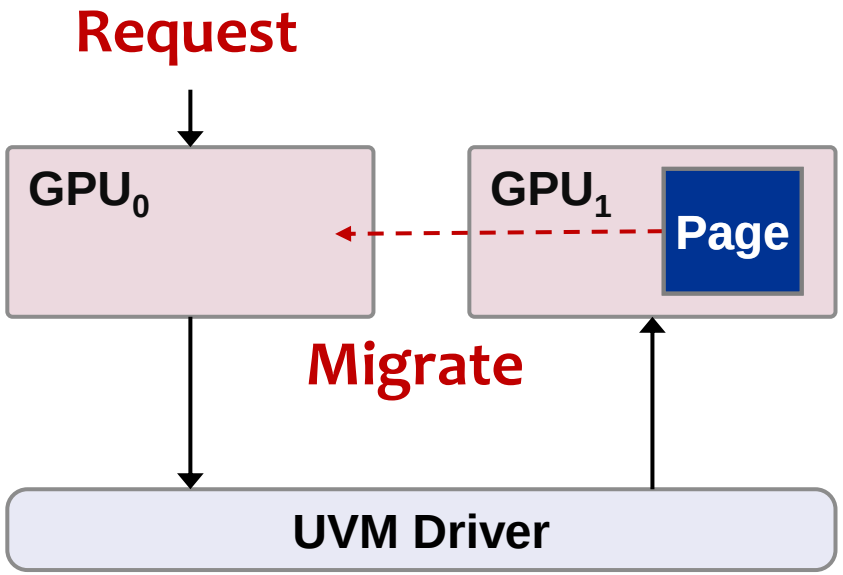

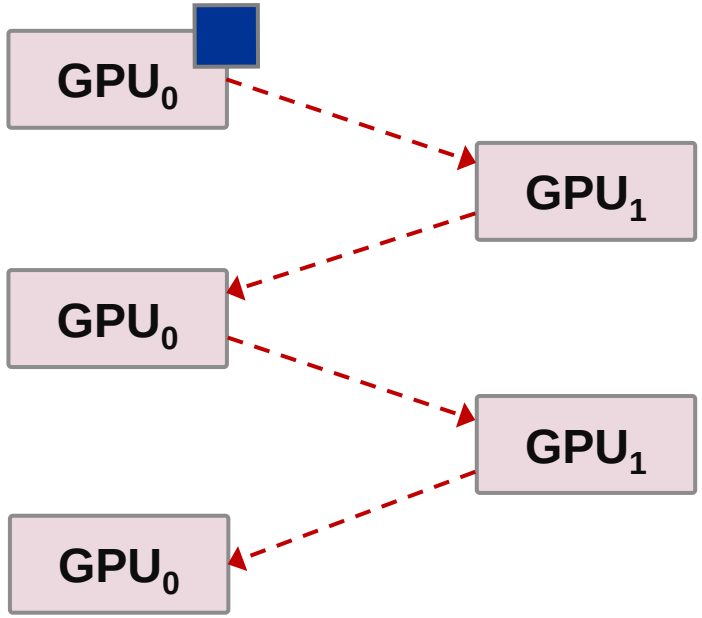

Every memory access is local

Ping-Pong effect

Multi-GPU Page Placement Schemes

2. Access Counter-based Migration: Reaching threshold, Migrate.

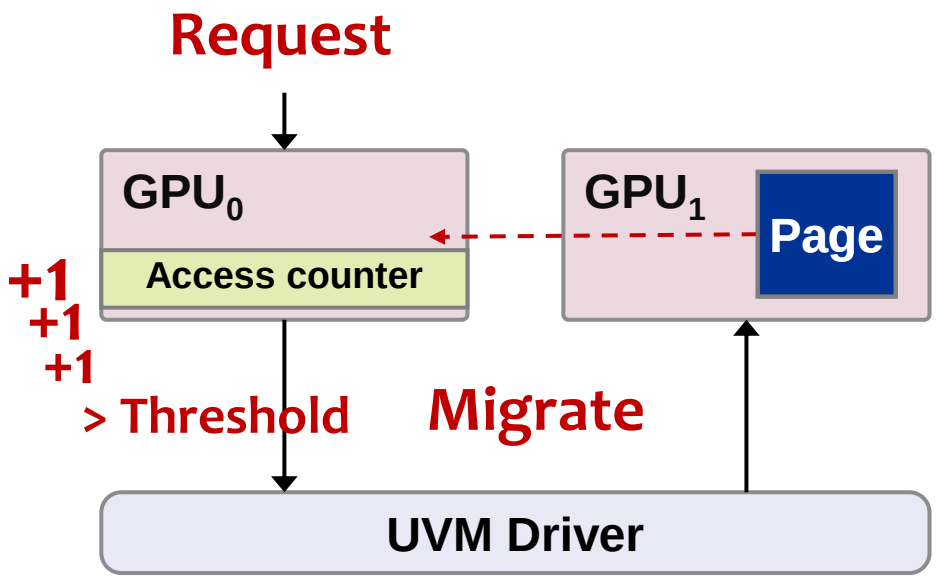

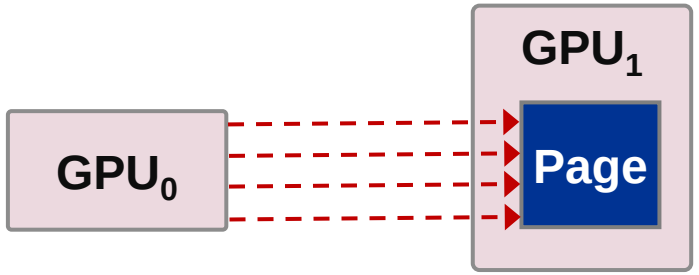

Reduce Ping-Pong migration

Substantial remote access

Multi-GPU Page Placement Schemes

3. Page Duplication: Read, Duplicate.

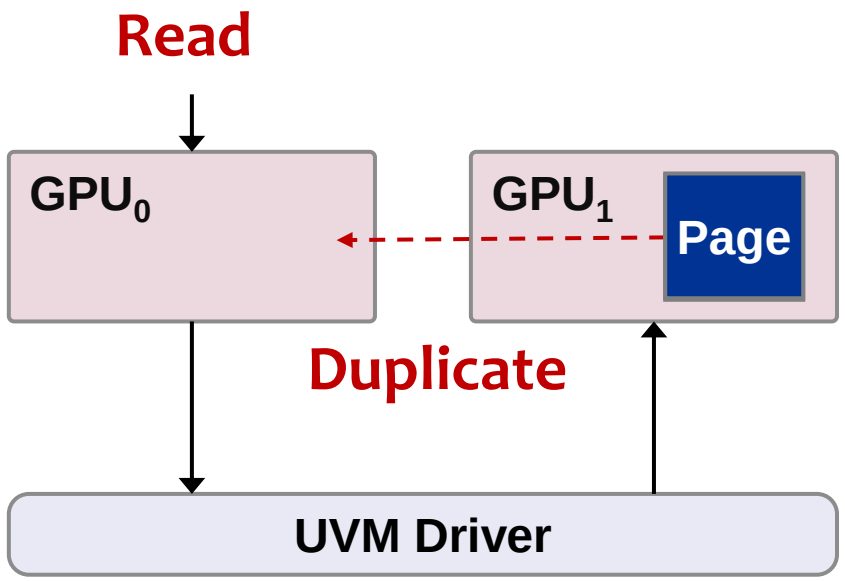

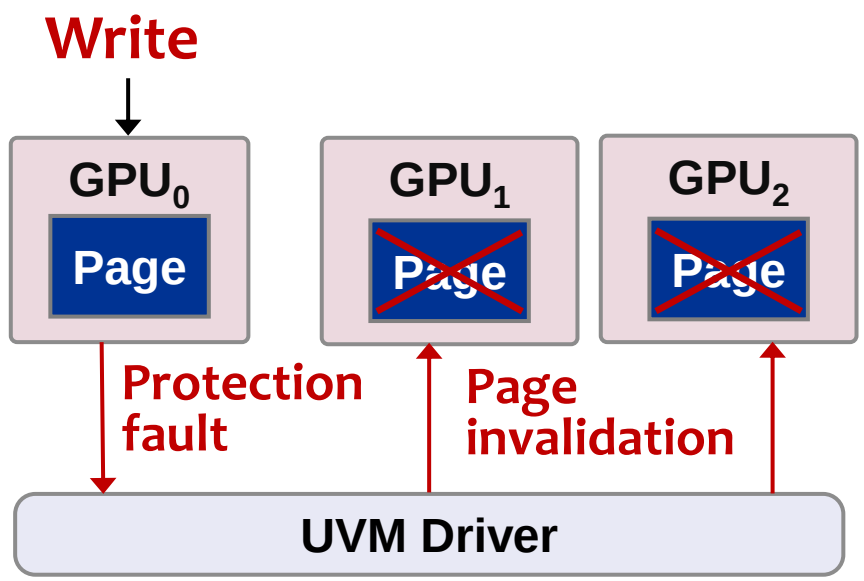

Reduce migration & remote access

Significant collapsing overhead

Performance of Page Placement Schemes

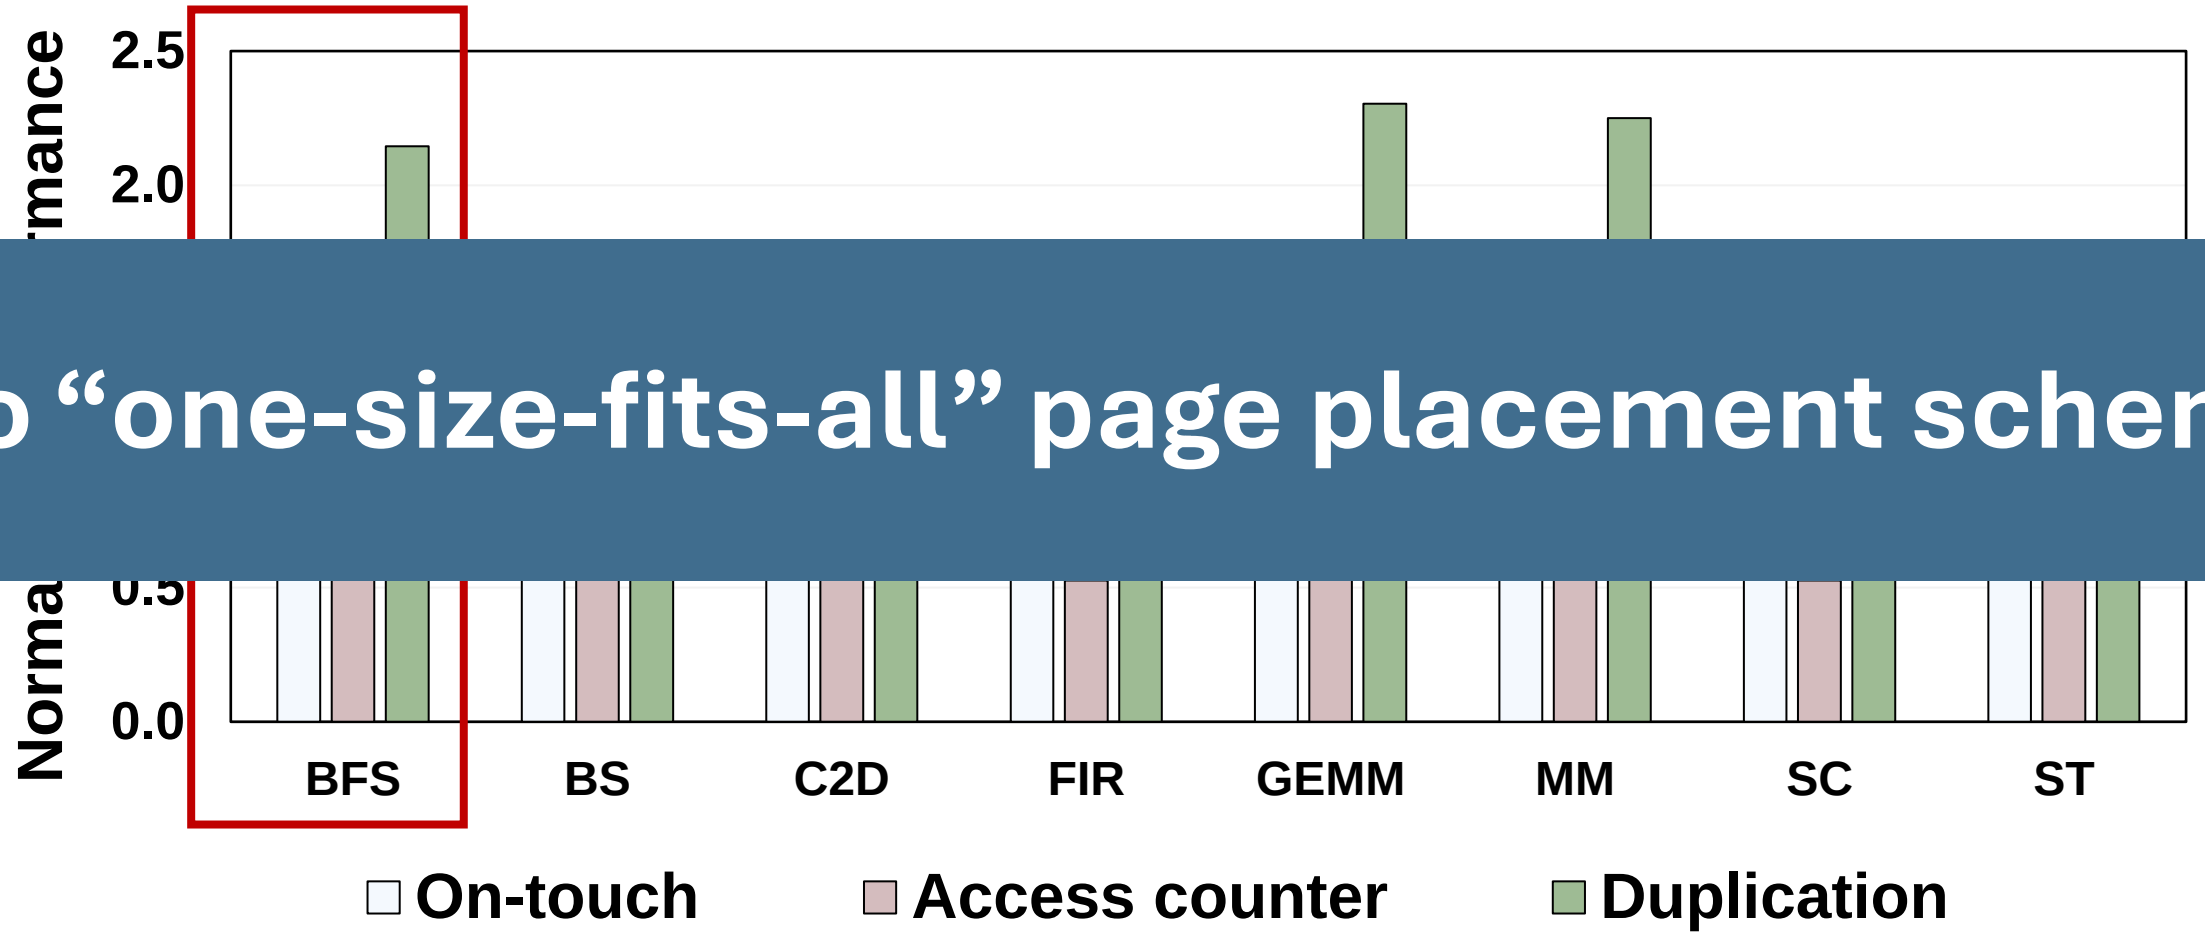

No “one-size-fits-all” page placement scheme

Page Access Characteristics

- The page-sharing / read-write patterns vary **within the same application**

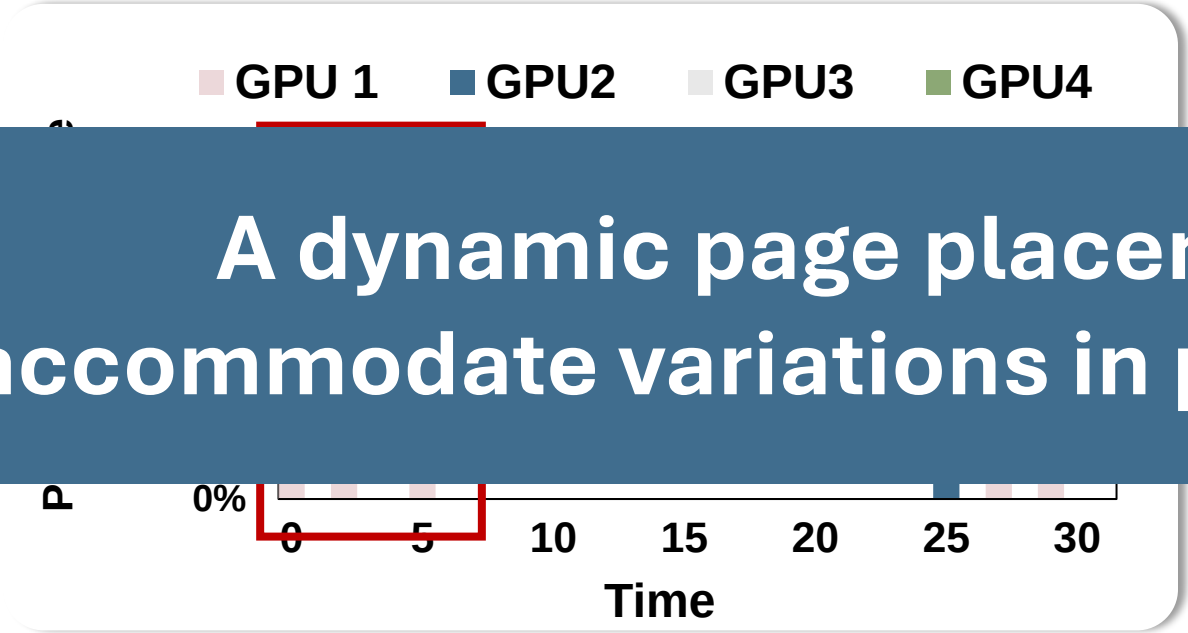

Private / Share

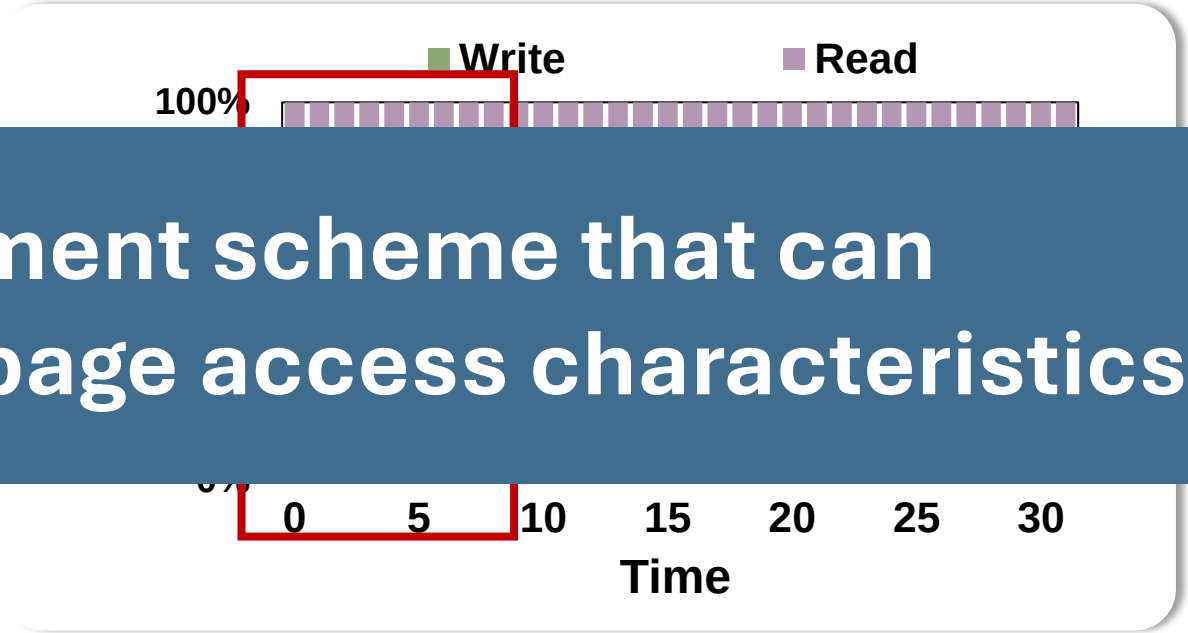

Read / Write

A dynamic page placement scheme that can accommodate variations in page access characteristics

Problem Summary

Problem:

Delivered performance is constrained by NUMA overhead

No "one-size-fits-all" page placement scheme

Goal:

Effectively reduce NUMA overhead in multi-GPU by determining page placement scheme at runtime

GRIT (Fine-**GR**ained dynam**I**c page placemen**T**)

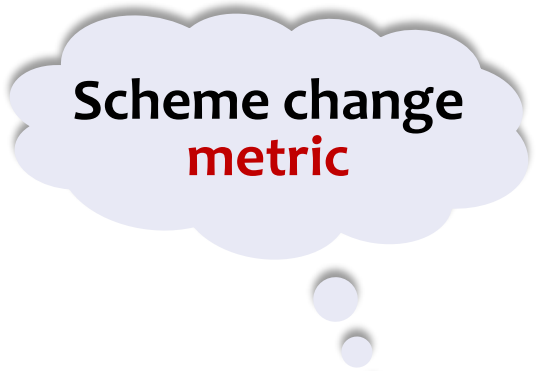

Dynamically determine page placement scheme at runtime

GRIT – Scheme Change Metric

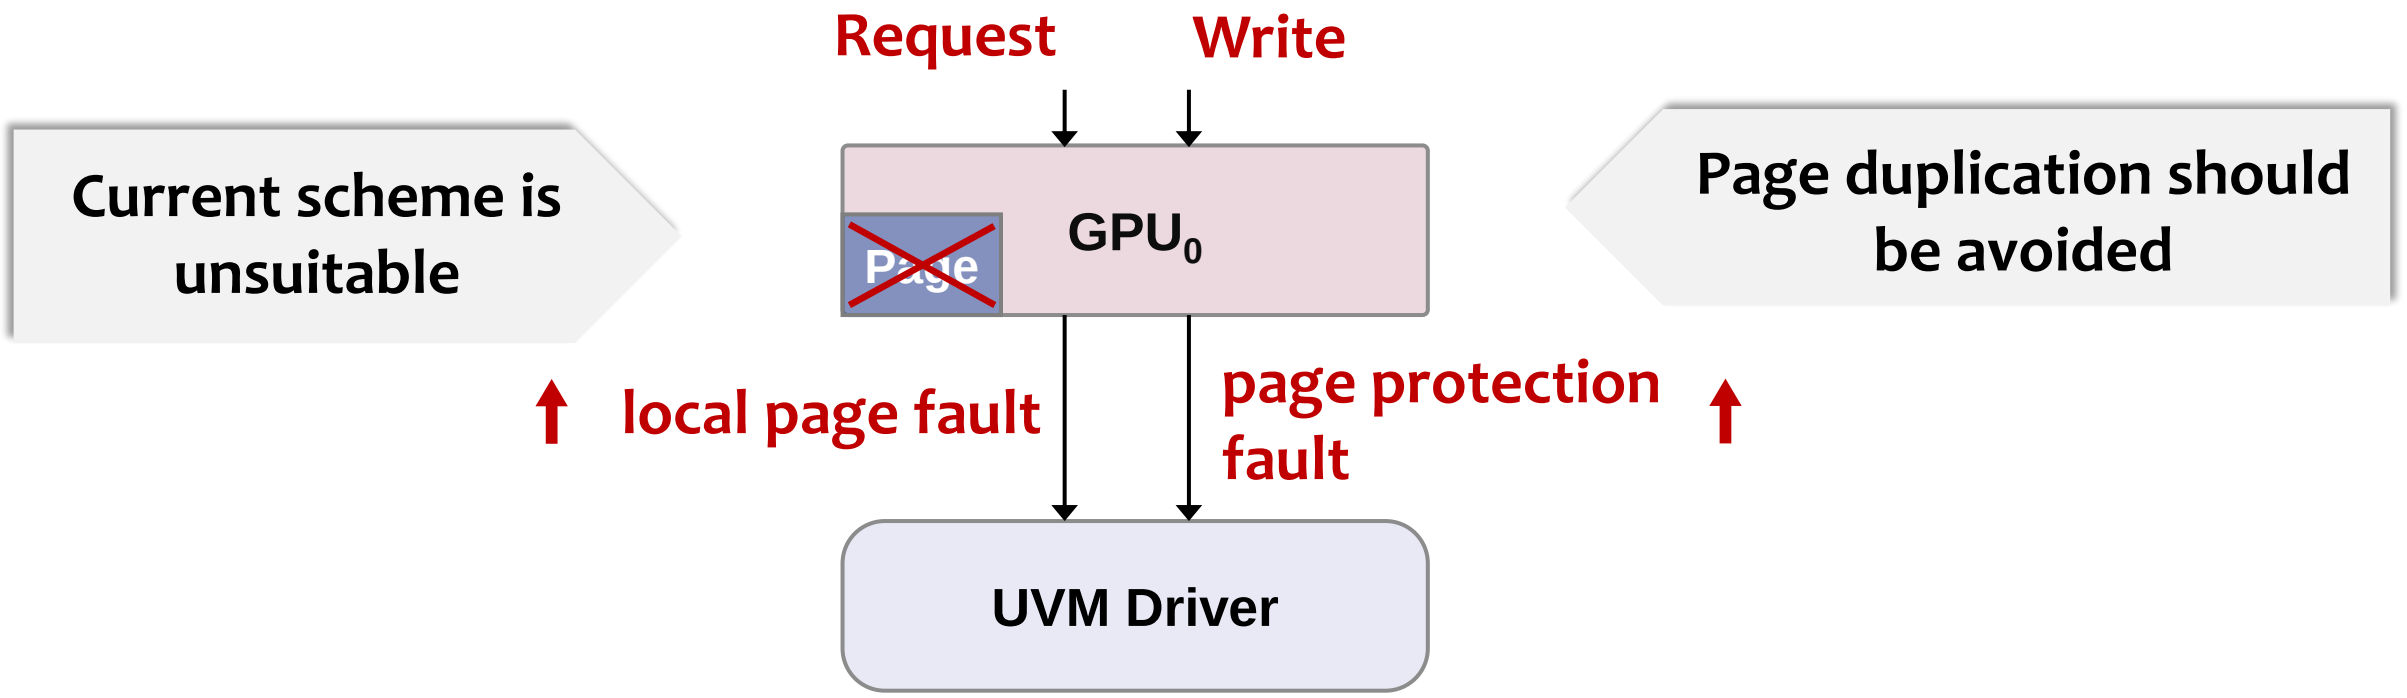

Indicator: Number of page faults (local page fault & page protection fault)

GRIT: Dynamic Page Placement Scheme

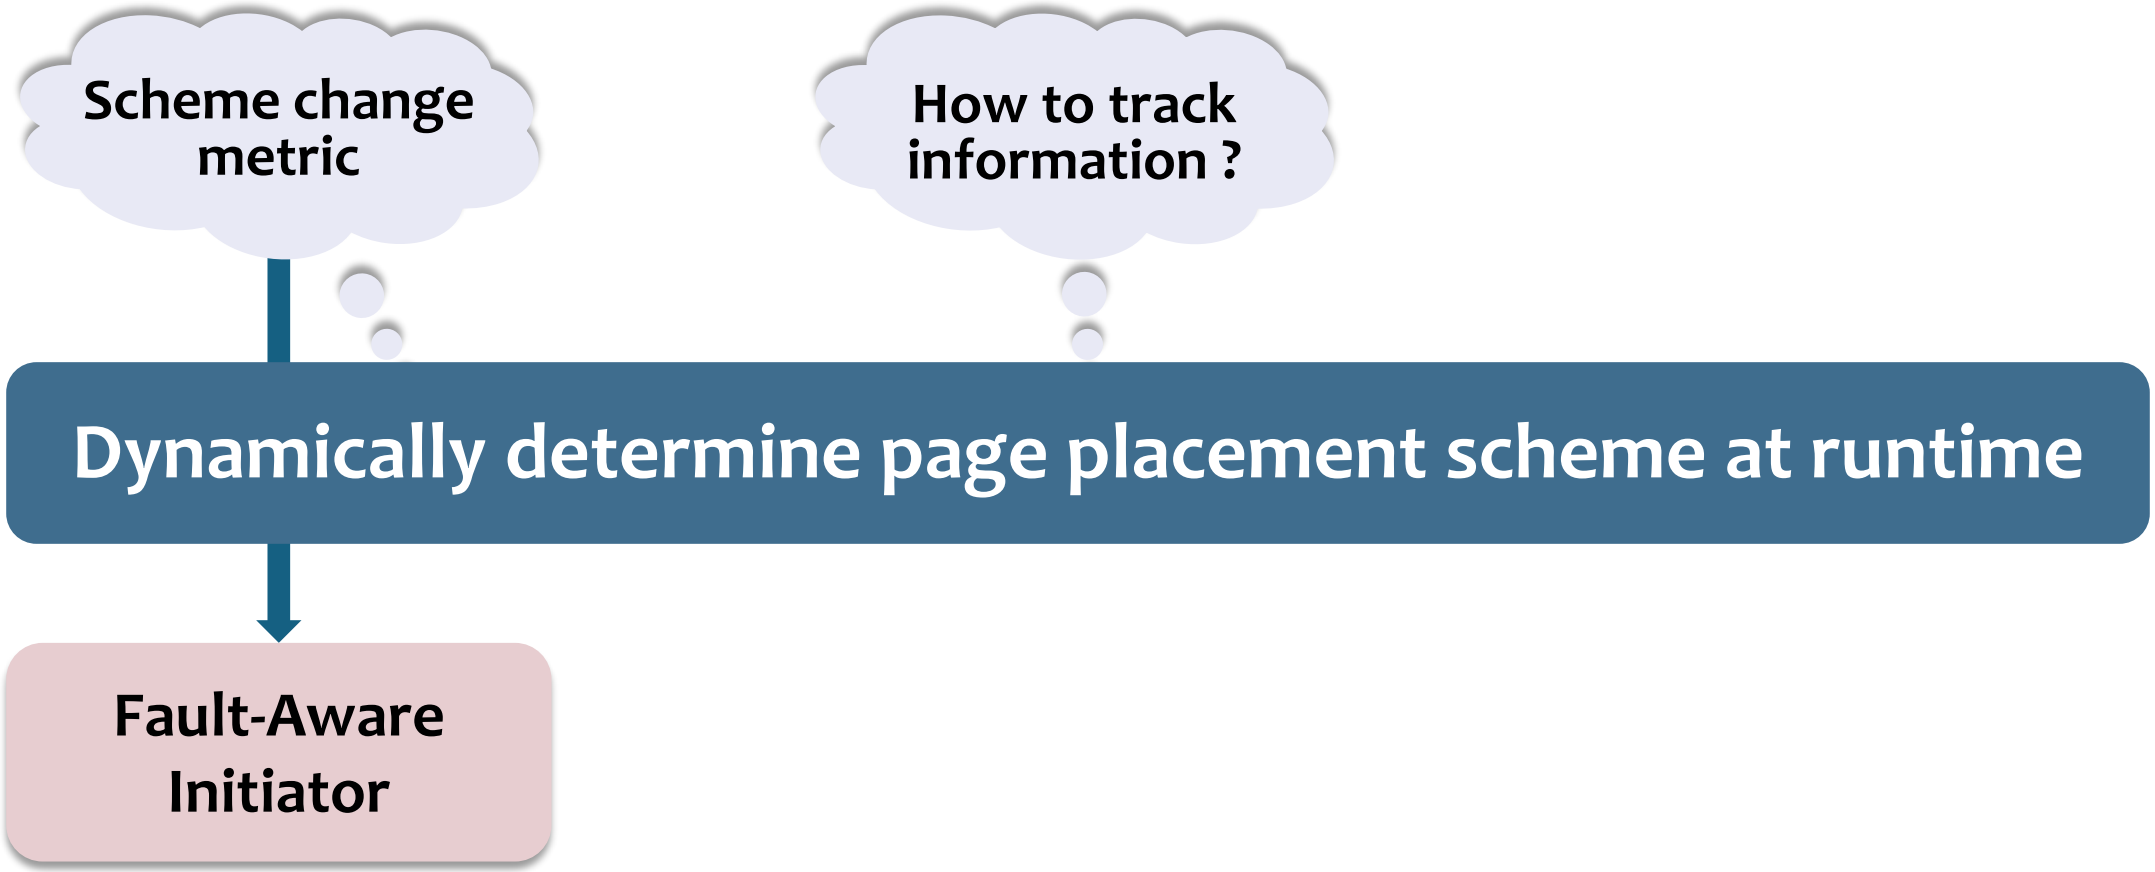

GRIT – Page Attribute Table and Cache (PA-Table & Cache)

Additional memory access

Facilitate lookup

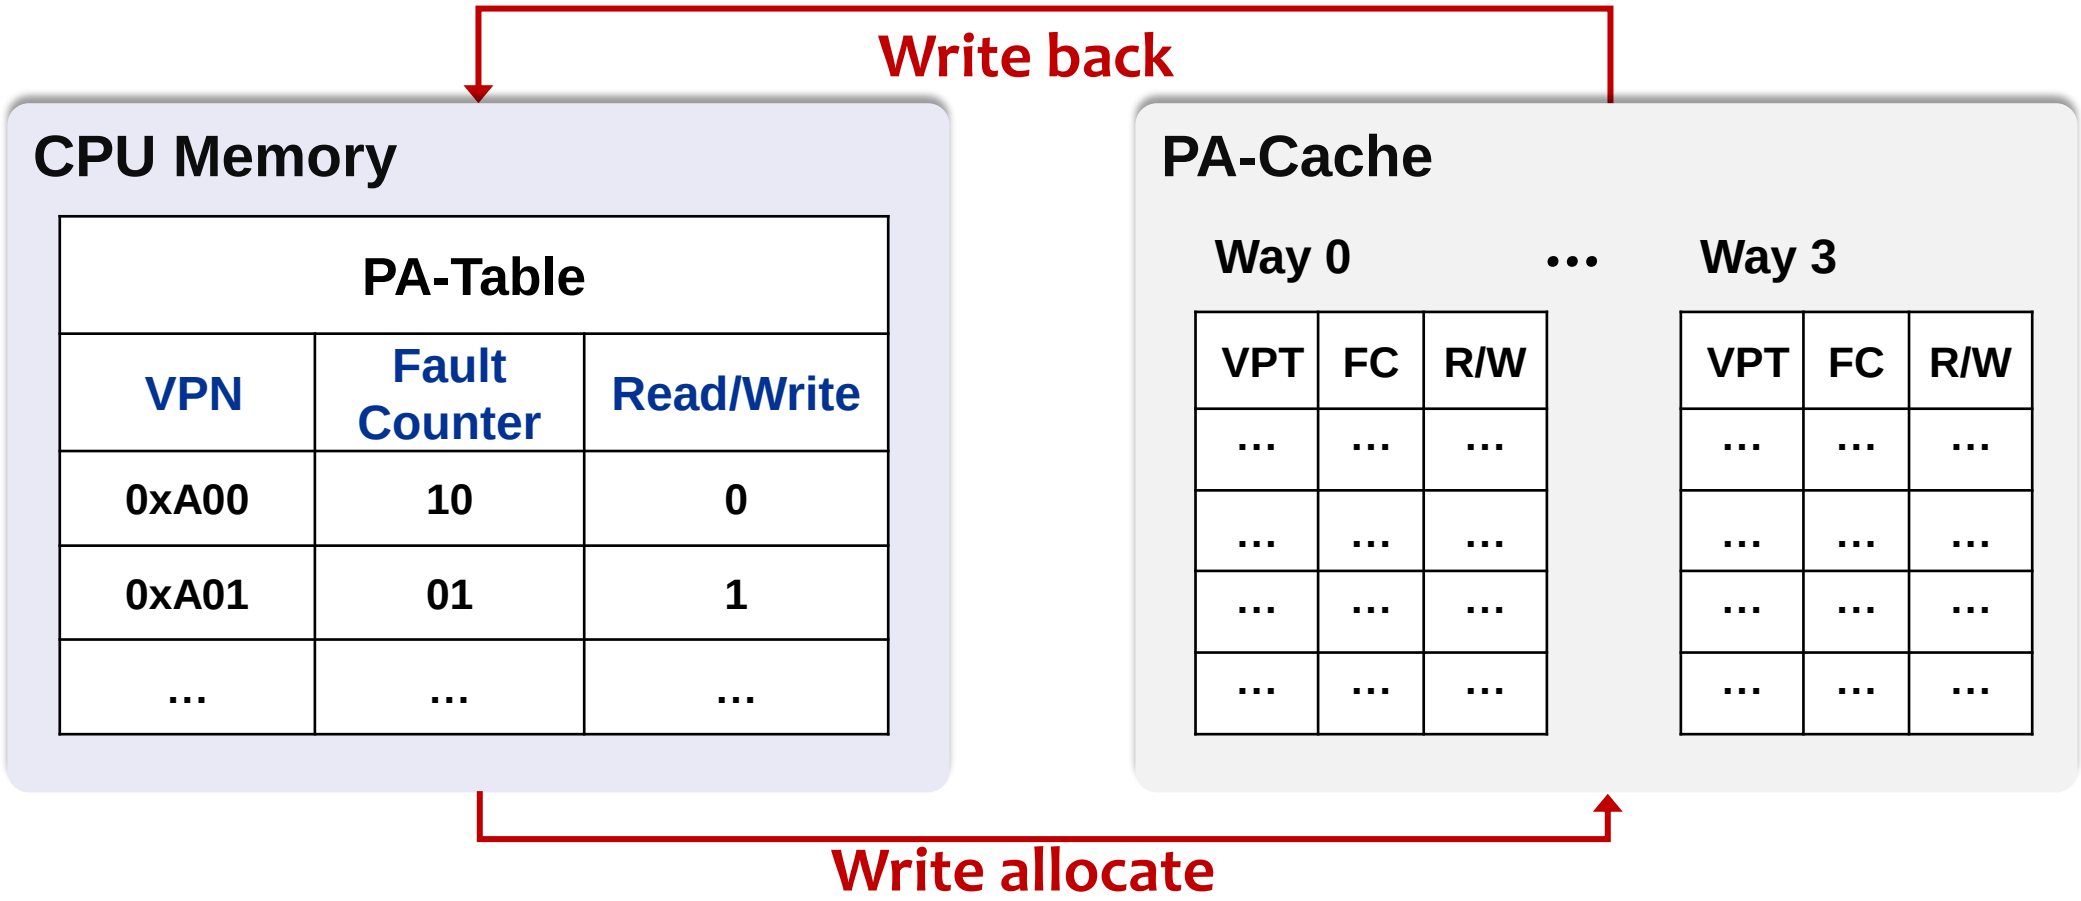

GRIT: Dynamic Page Placement Scheme

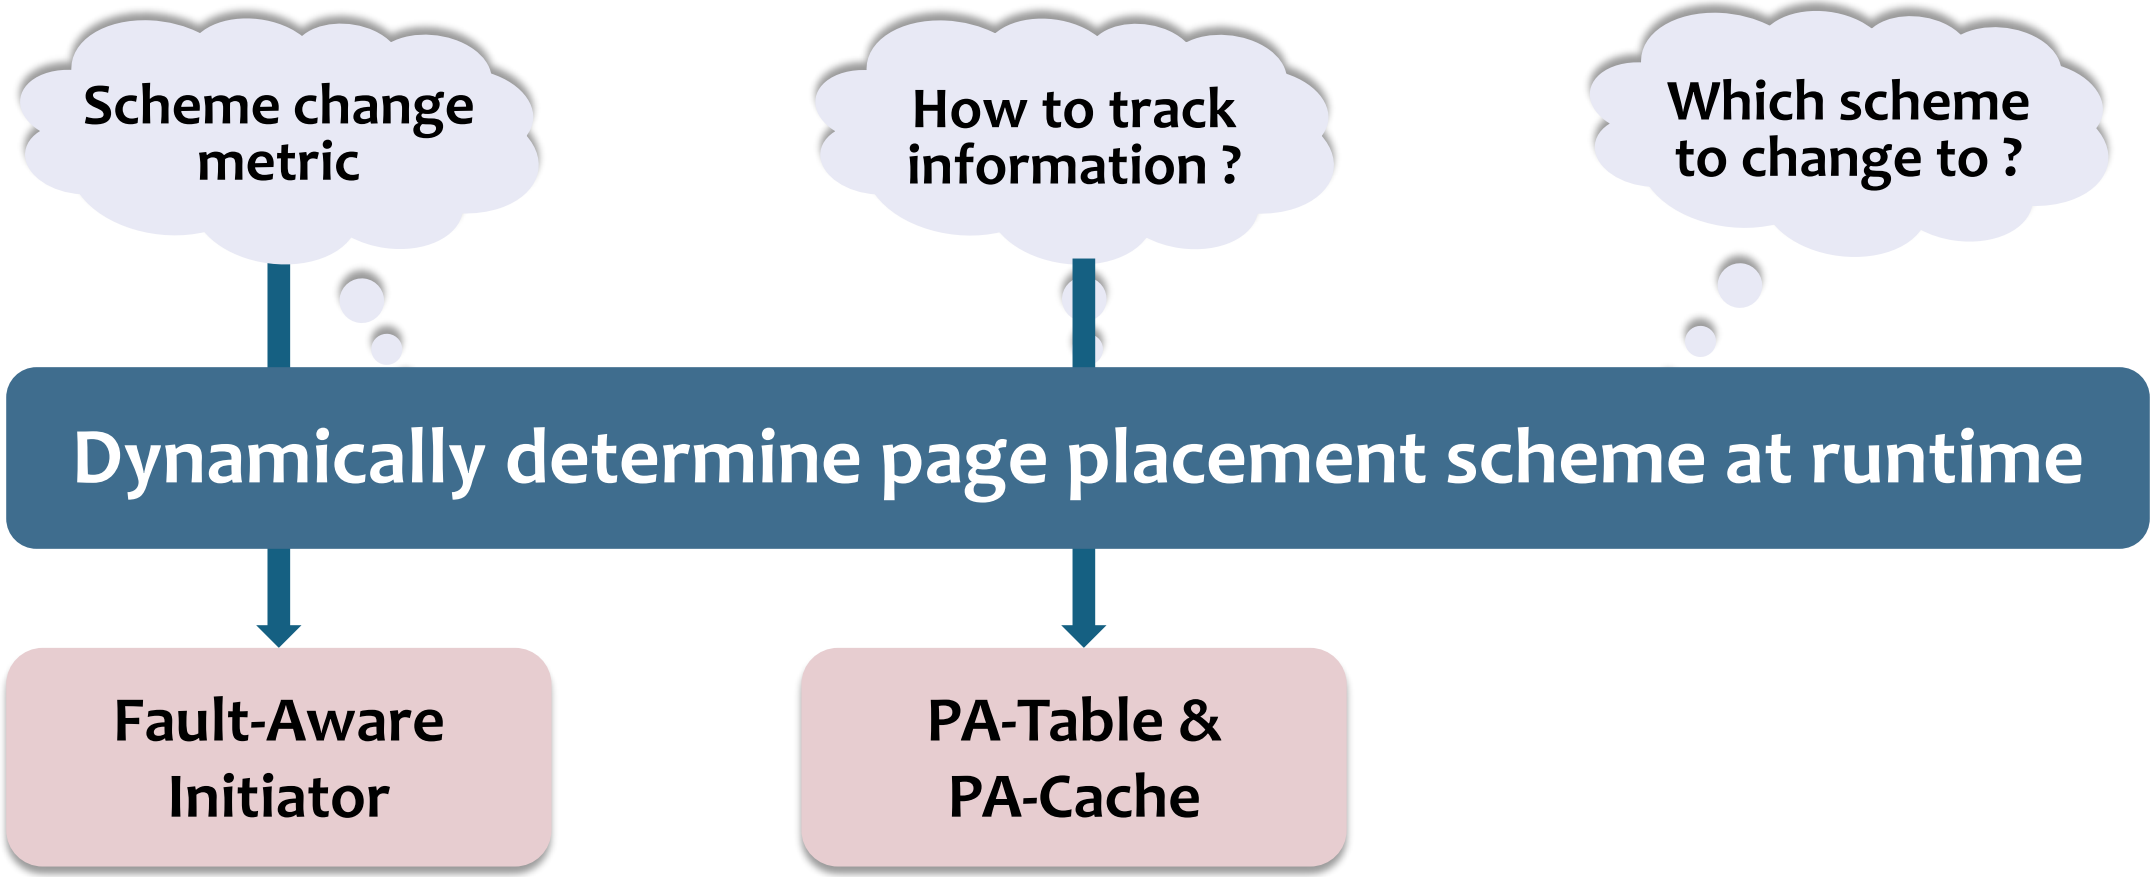

GRIT – Which Scheme

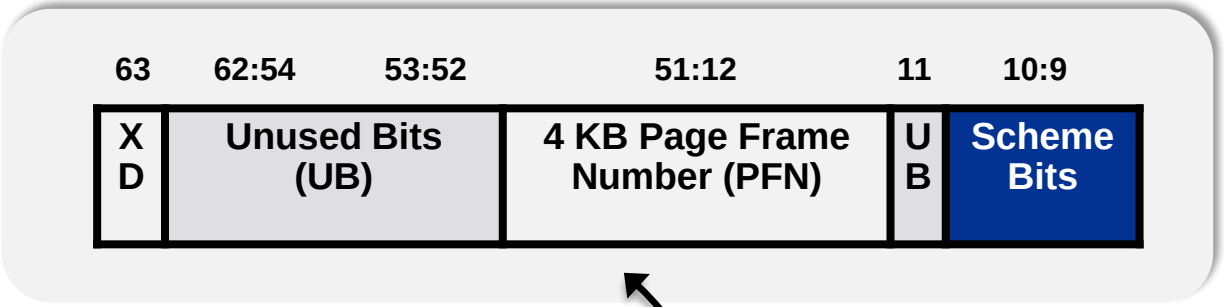

On-touch

Get access information from PA-Table /PA-Cache

FC Reach threshold?

True

Change Scheme

All read?

True

Duplication

False

Access-counter

Scheme information is stored in host PTE.

GRIT: Dynamic Page Placement Scheme

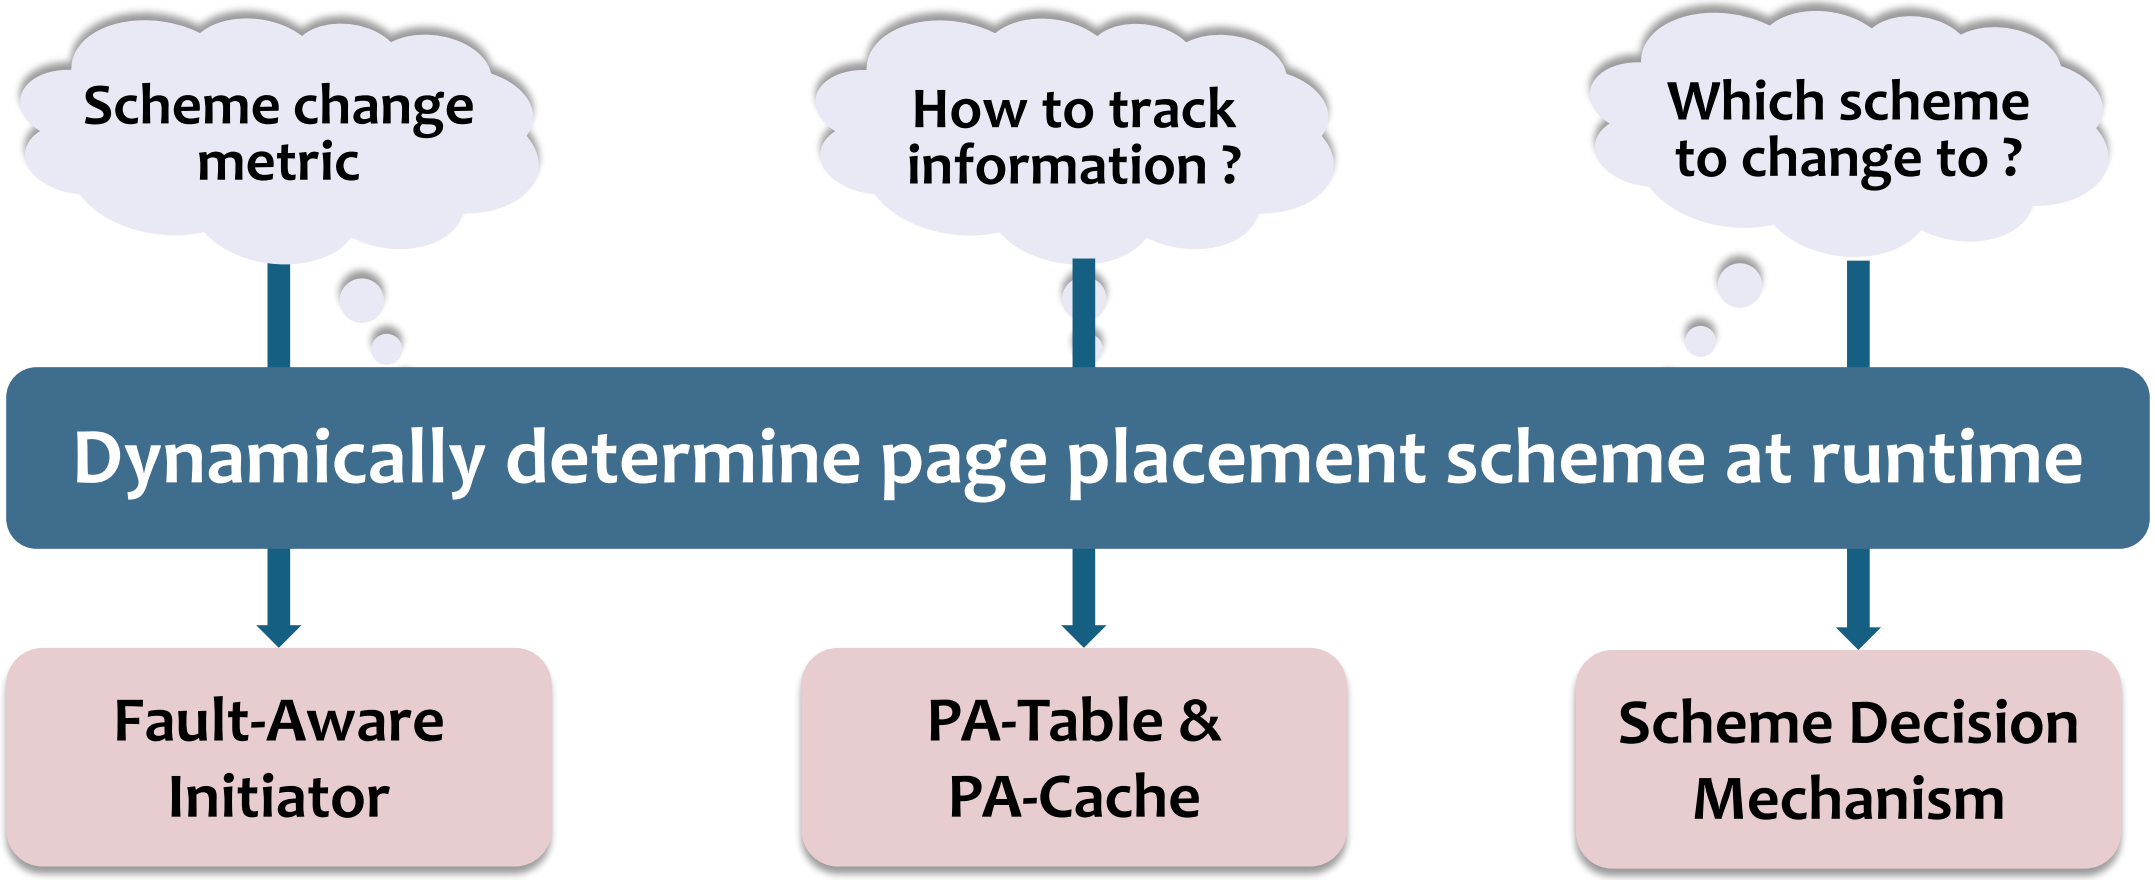

GRIT: Dynamic Page Placement Scheme

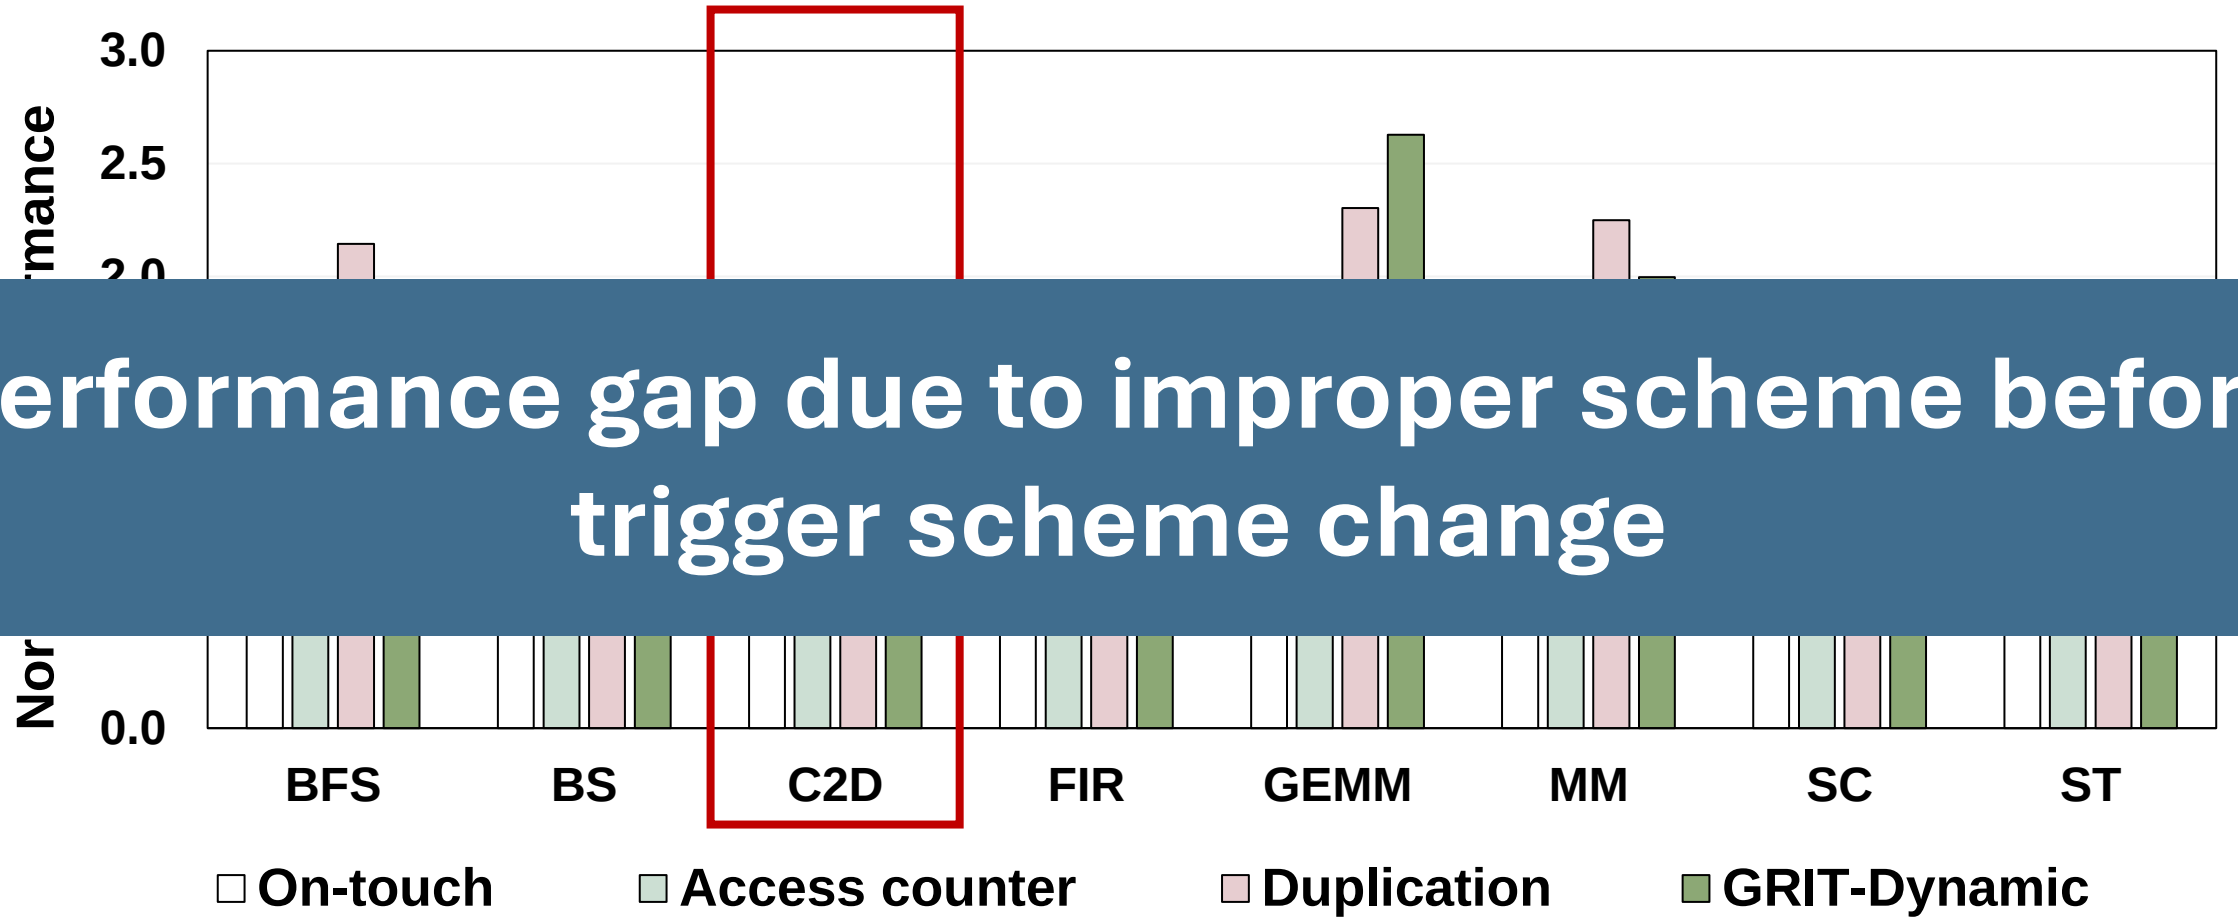

Performance gap due to improper scheme before trigger scheme change

Page Attributes Characterization

- The neighboring pages tend to exhibit similar access attributes

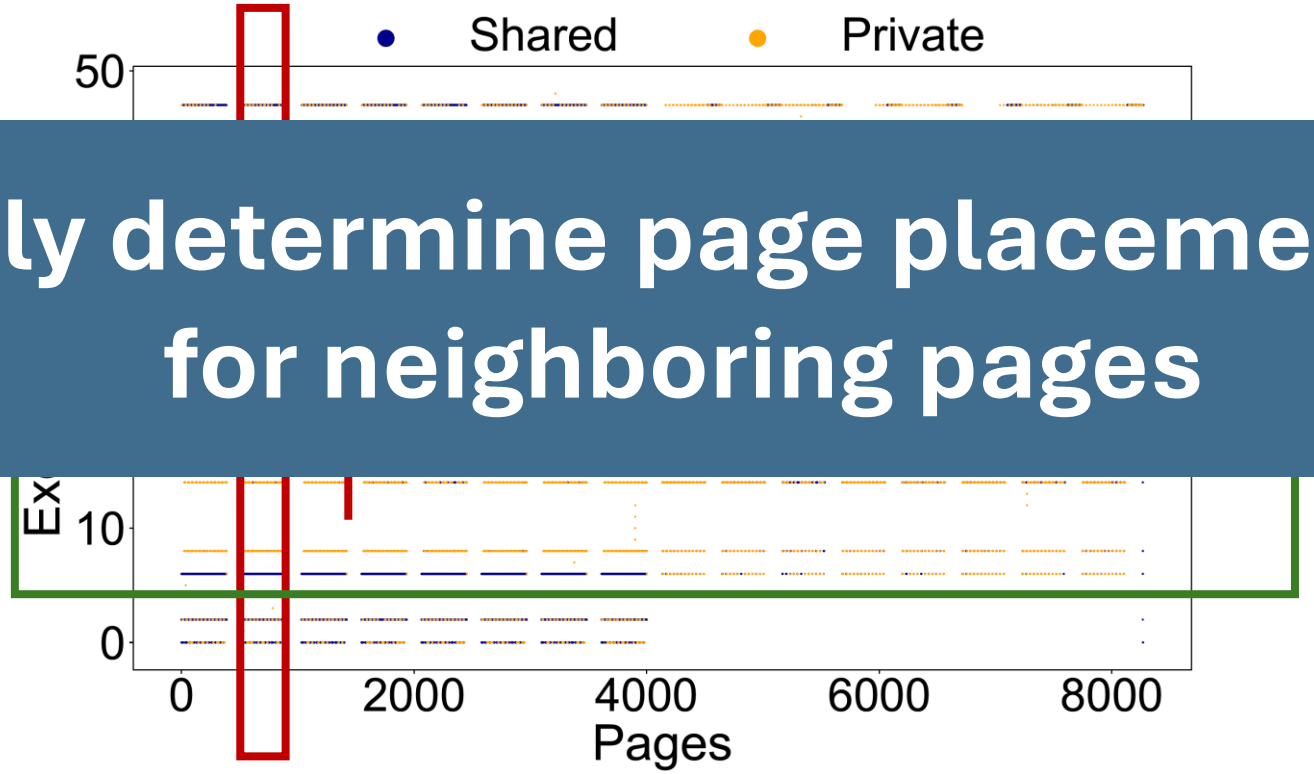

Proactively determine page placement scheme for neighboring pages

GRIT: Neighboring-Aware Prediction

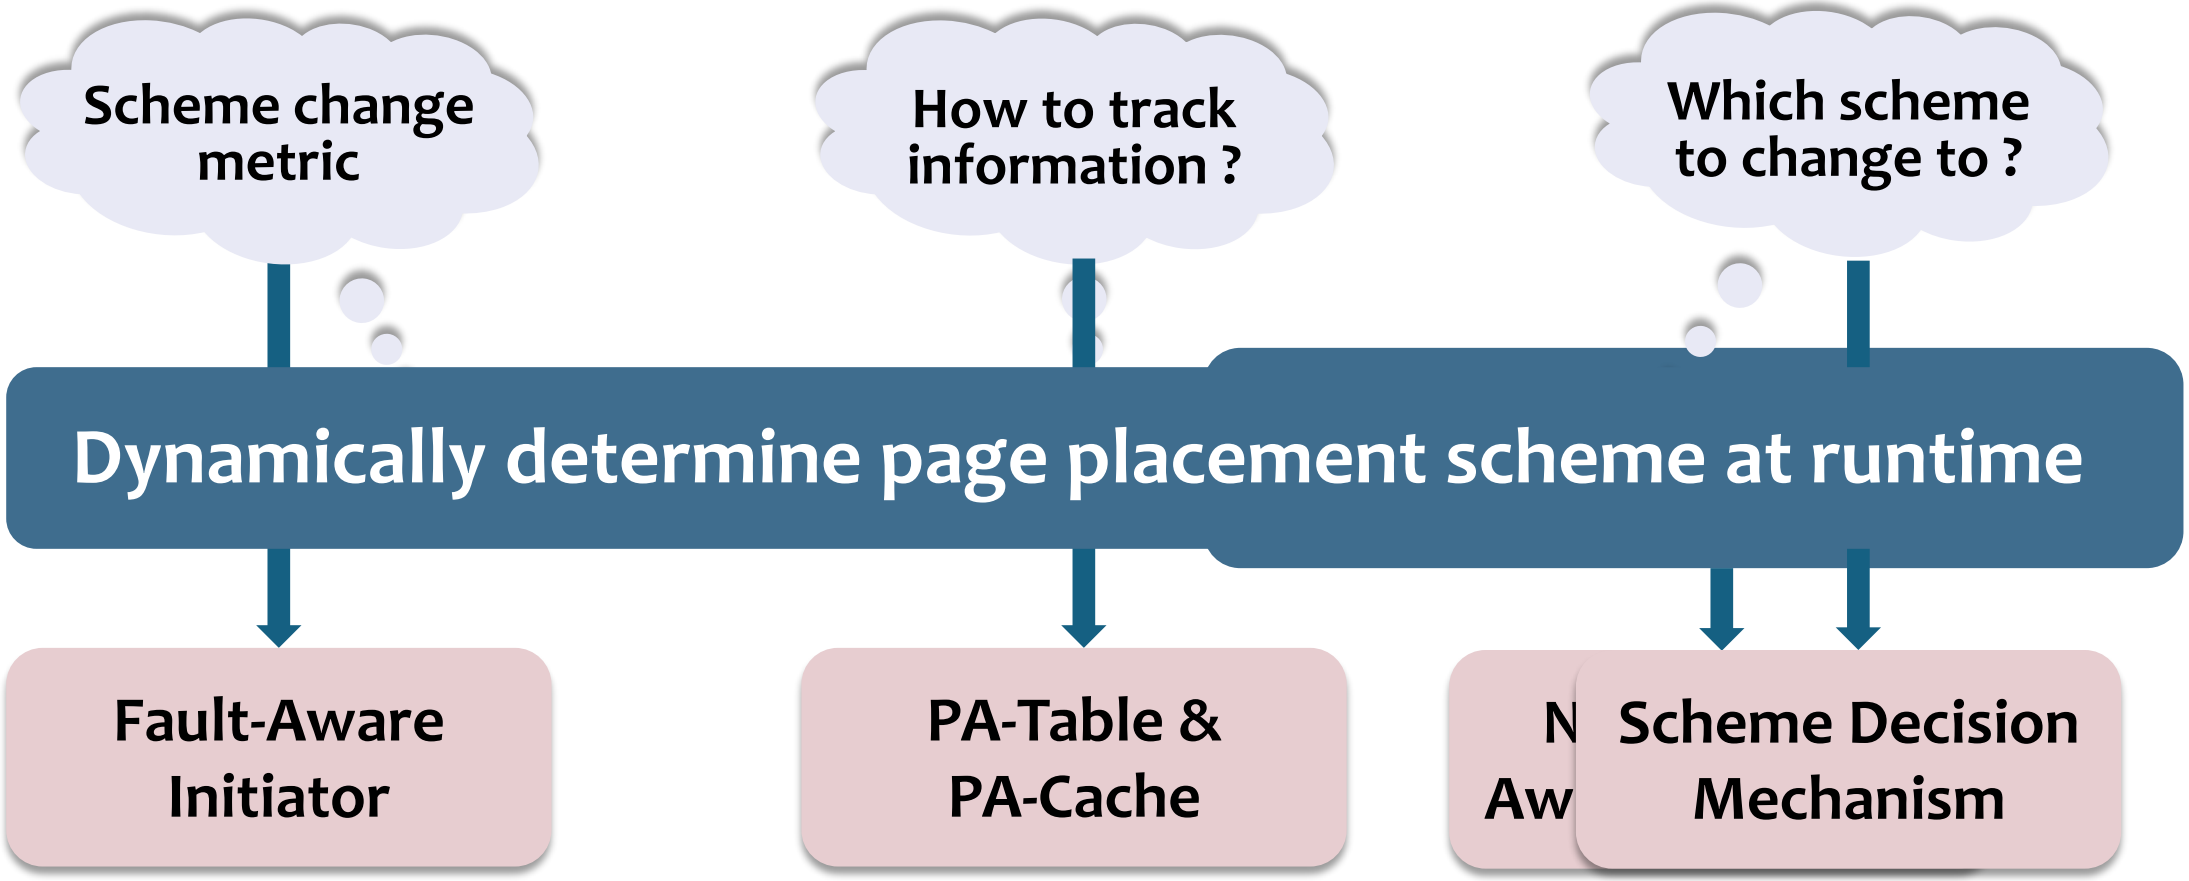

GRIT: Neighboring-Aware Prediction

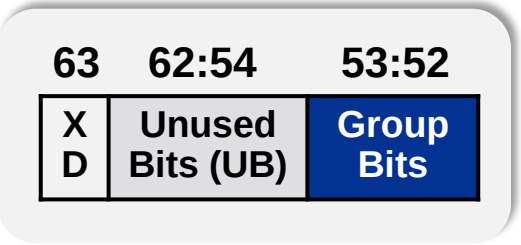

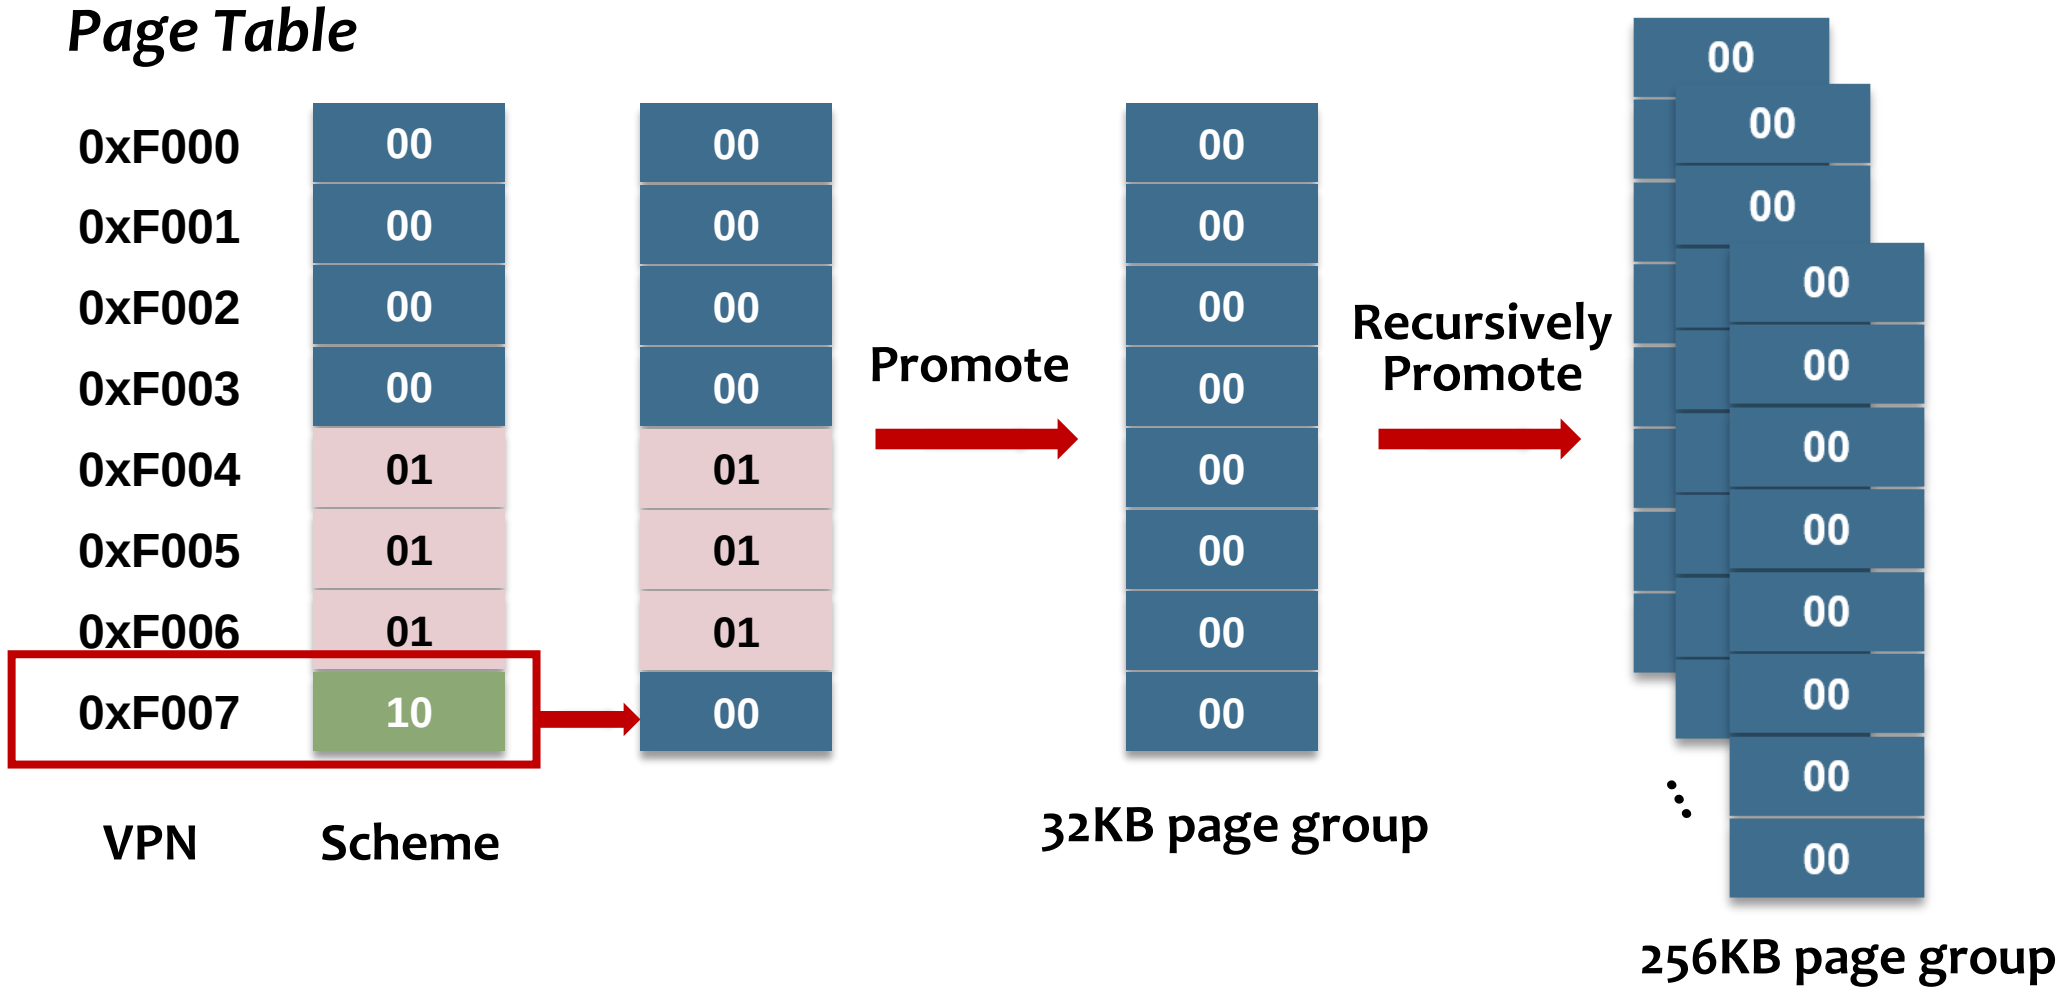

GRIT – Put All Together

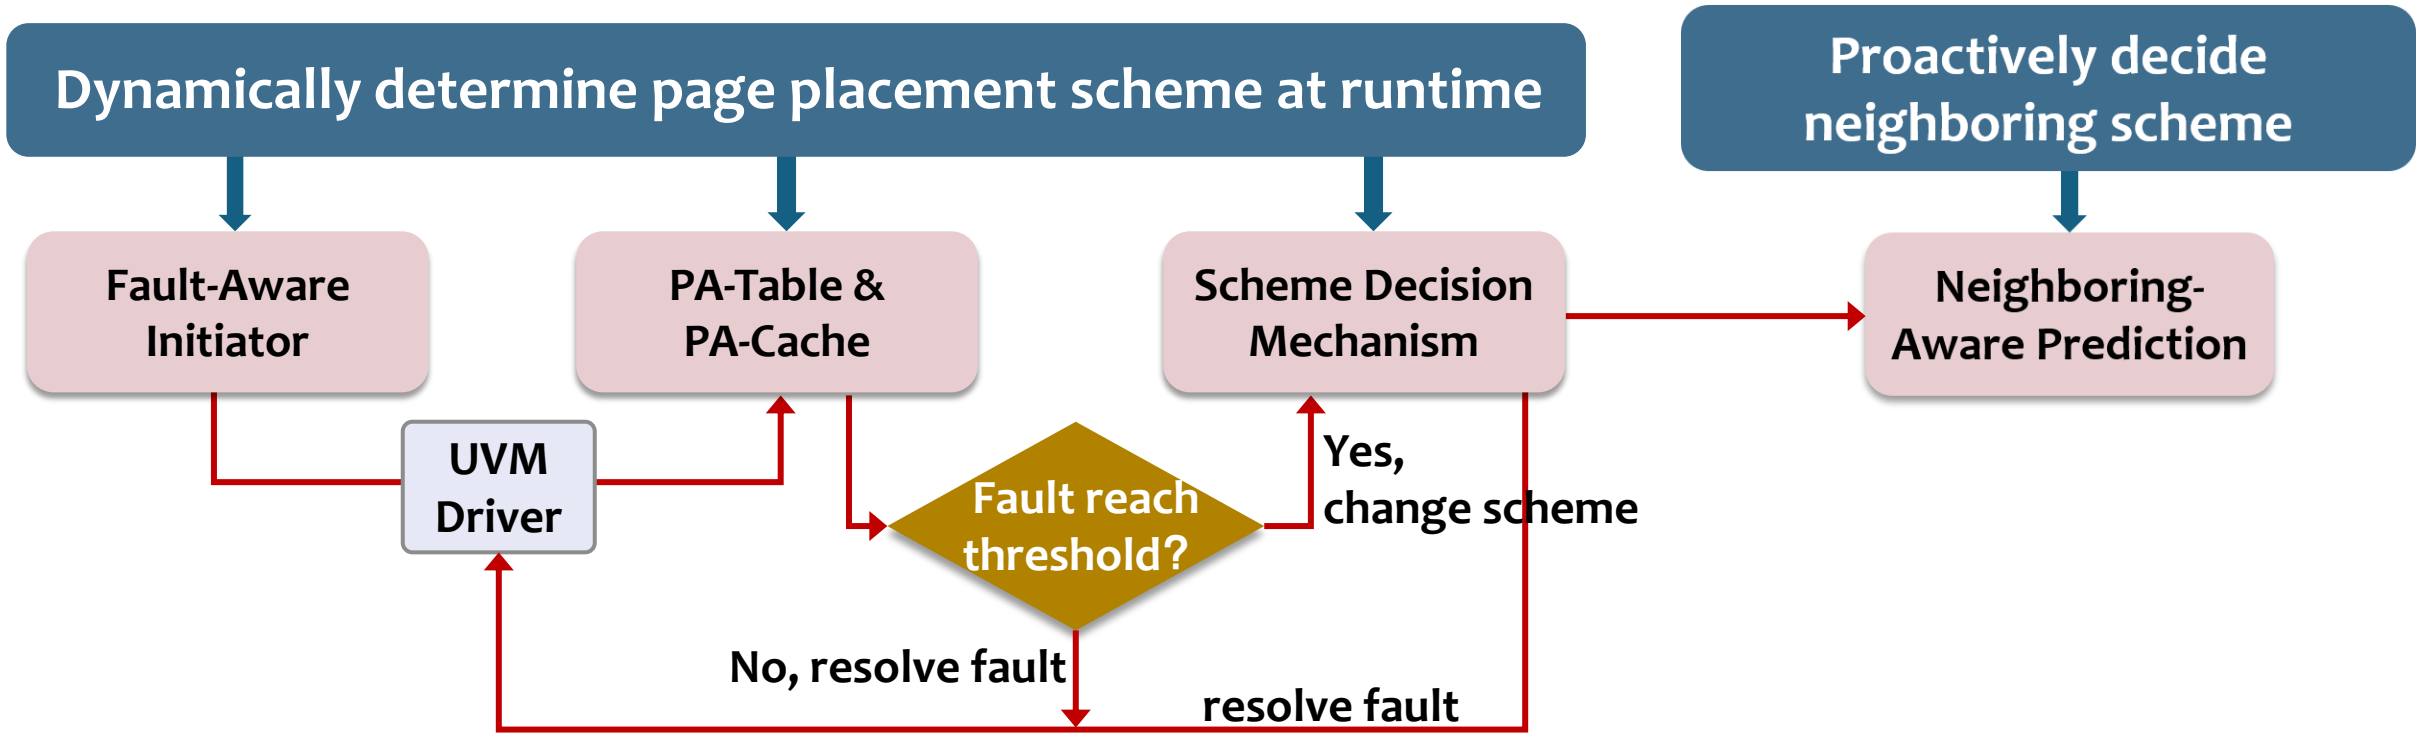

Methodology

- Simulator: MGPUSim [ISCA 19’]
- Workloads: 8 applications from Hetero-Mark, AMDAPPSDK , SHOC, and DNN Mark benchmark suites, including random, adjacent, and scatter-gather access patterns.

Detailed page placement scheme modeling in paper

Evaluation – Overall Performance

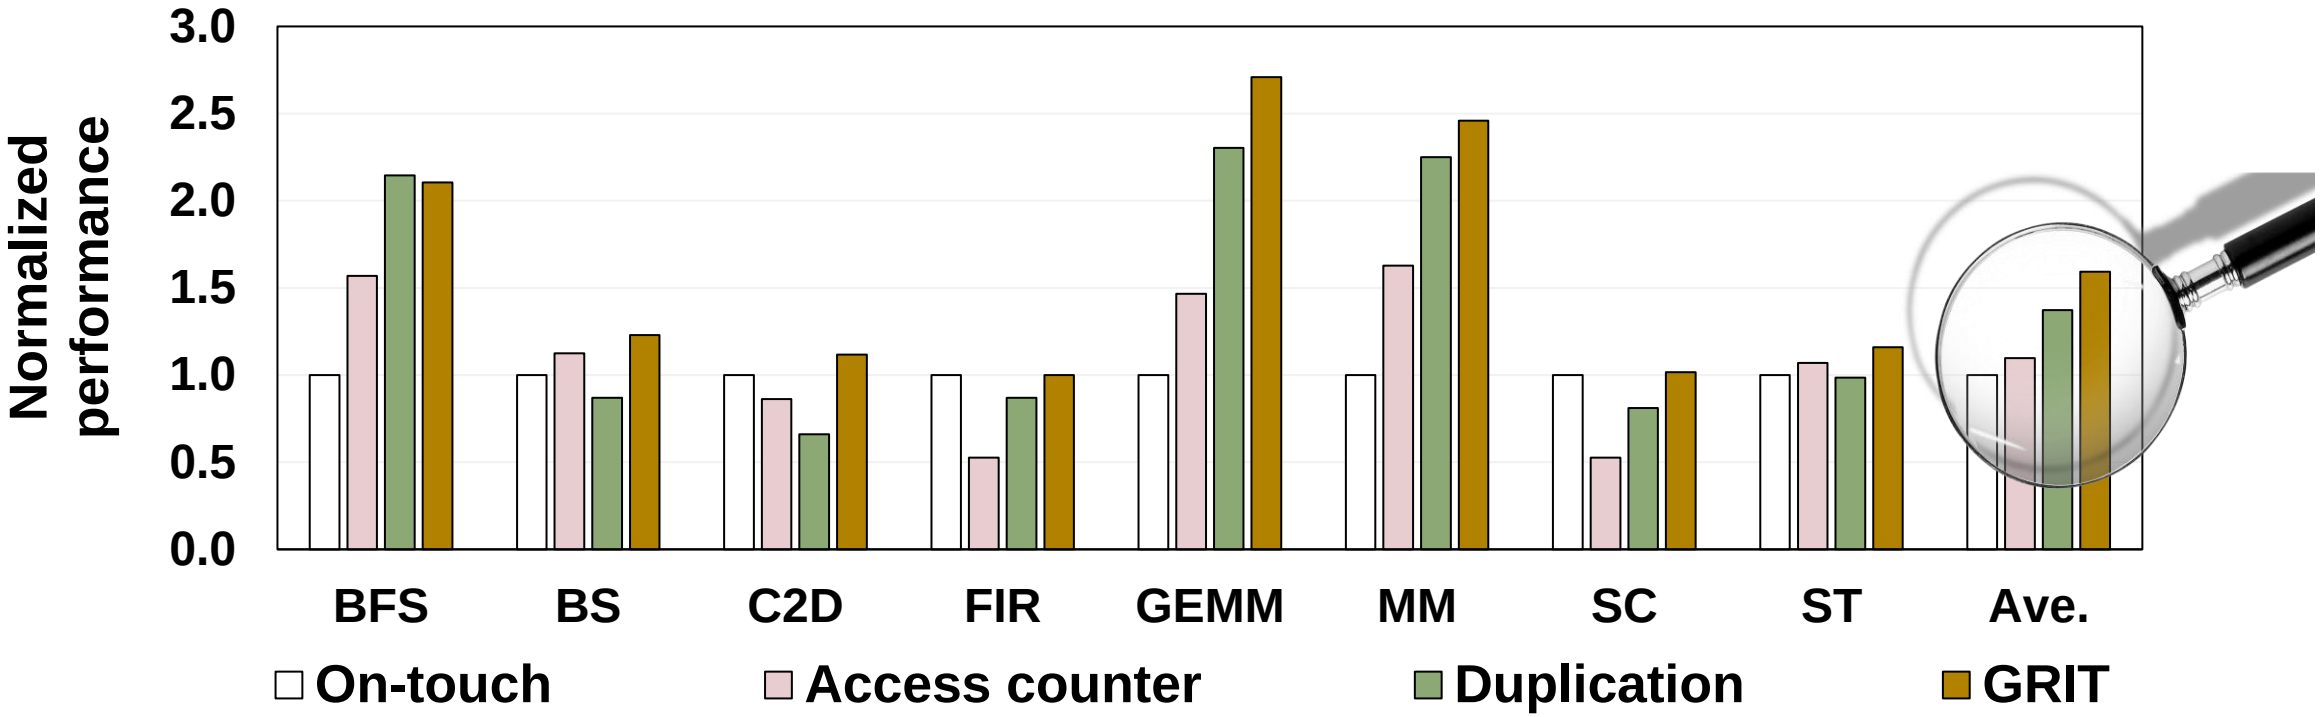

GRIT achieves *60%, 49%, and 29%* performance improvement compared to uniformly employing on-touch, access counter-based, and page duplication scheme.

University of Pittsburgh

Evaluation – Scheme Breakdown

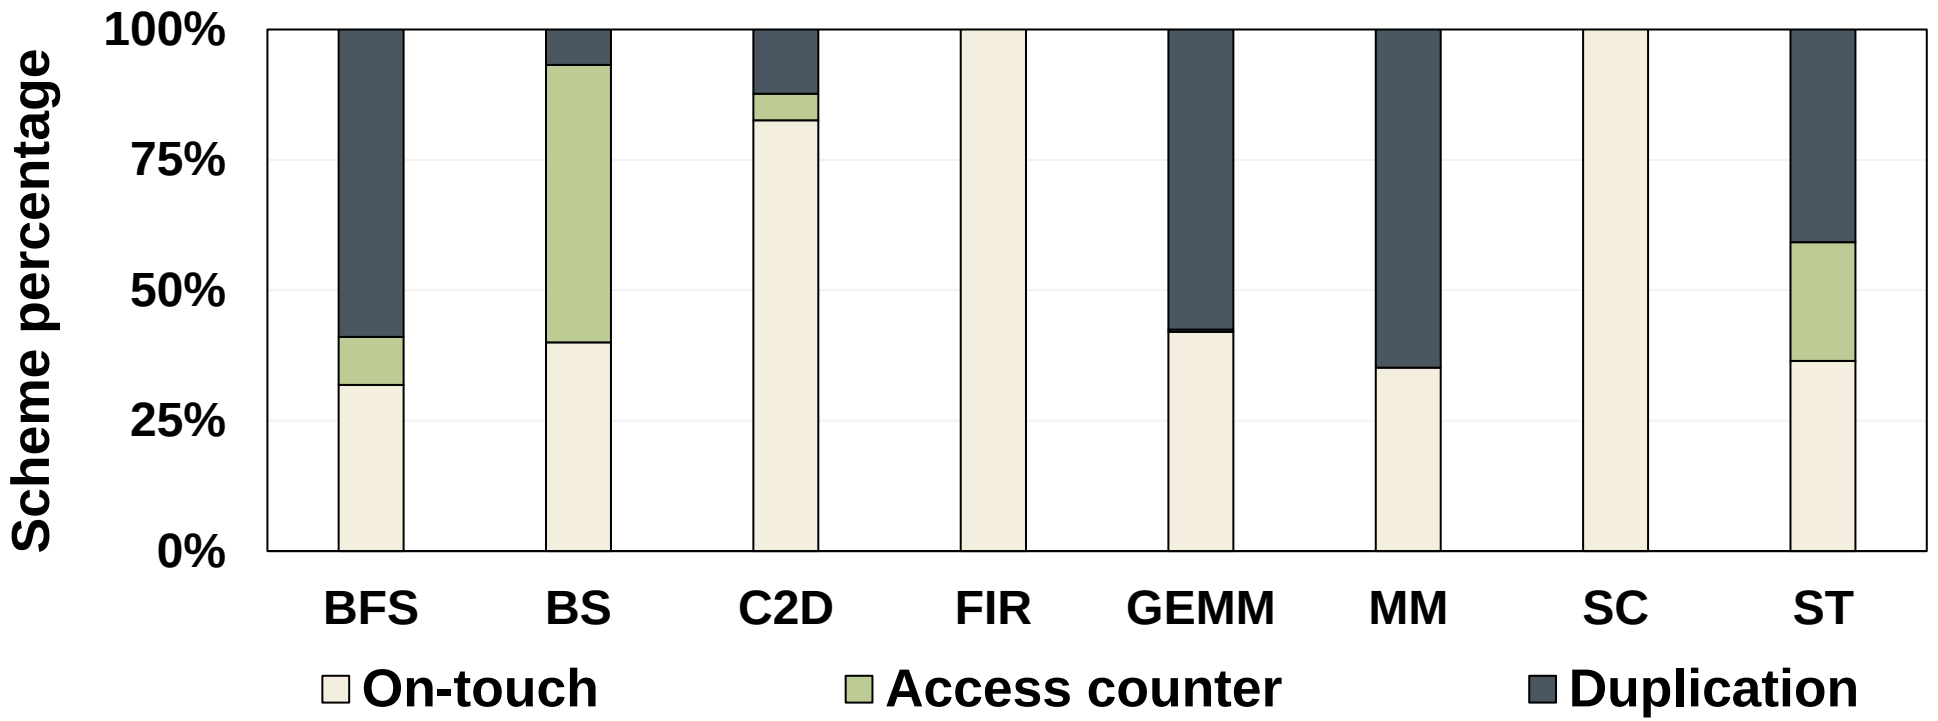

GRIT is able to distinguish page attributes and consistently select the most suitable scheme accordingly.

Summary

Problem: NUMA overheads in multi-GPU systems

- *No "**one-size-fits-all**" page placement scheme*

GRIT:

A. Dynamic page placement scheme **determines schemes in a fine-grained manner**
B. Neighboring-aware prediction **proactively determines adjacent page scheme**

Improves **performance** by **60%** on average.

University of Pittsburgh

Thanks! Q&A

GRIT: Enhancing Multi-GPU Performance with Fine-Grained Dynamic Page Placement

Yueqi Wang*[1], ***Bingyao Li****[1], Aamer Jaleel[2], Jun Yang[1], Xulong Tang[1]

[1]University of Pittsburgh, [2]NVIDIA